

## Mini Project #1

## UT dataset

- ```
princebirring1992@ubuntu-14:~/Deep-Learning/Torch_$ cd Mini_Project/
princebirring1992@ubuntu-14:~/Deep-Learning/Torch_/Mini_Project$ ls
Readme.md  train_mnist.lua  train_mnist_Remove_Bug.lua
princebirring1992@ubuntu-14:~/Deep-Learning/Torch_/Mini_Project$ th

  ____  _
 / ___|| | | |
| |___| |_| |
 \___ \|  __/
      | | | |
      |_|_|_|

| Torch7
| Scientific computing for Lua.
| Type ? for help
| https://github.com/torch
| http://torch.ch

th> require'dp'
{
  XpLog : {...}
  mkdir : function: 0x41f6f338
  FKDKaggle : {...}
  SAVE_DIR : "/home/princebirring1992/save"
  ListView : {...}
  is file : function: 0x41f6f5e8
  BillionWords : {...}
  ImageNet : {...}
  SequenceView : {...}
  SentenceSet : {...}
  TextSet : {...}
  Cifar100 : {...}
  NotMnist : {...}
  download_file : function: 0x41f6de28
  reverseDist : function: 0x41f6e520
  TranslatedMnist : {...}
  XpLogEntry : {...}
}
```

2. Downloads the `trains_mnist.lua` file from Github(See <https://github.com/amir-jafari/Deep-Learning/tree/master/Torch> ). It is based on the following website <https://github.com/nicholas-leonard/slides/blob/master/torch7.md>, which may want to review before proceeding.

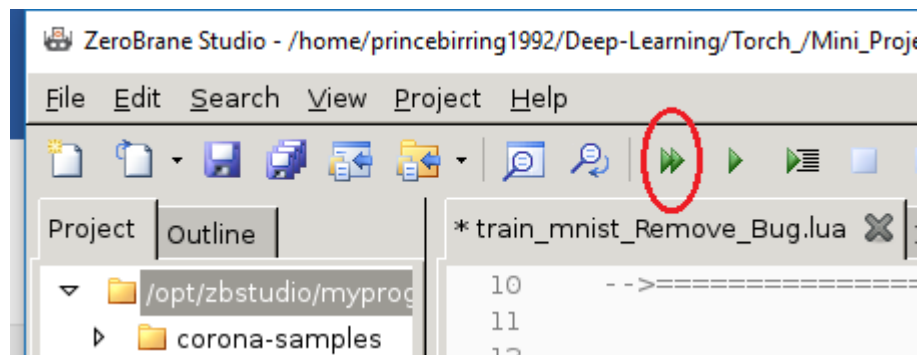
```
princebirring1992@ubuntu-14:~$ ls
caffe                                cudnn-8.0-linux-x64-v6.0.tgz      install-14-0
cuda                                cudnn-8.0-linux-x64-v6.0.tgz.1    log
cuda-repo-ubuntu1404-8-0-local-ga2_8.0.61-1_amd64.deb  data                               pycharm-comm
cuda-repo-ubuntu1404-8-0-local-ga2_8.0.61-1_amd64.deb.1  demos                             pycharm-comm
princebirring1992@ubuntu-14:~$ git clone https://github.com/amir-jafari/Deep-Learning.git
Cloning into 'Deep-Learning'...
remote: Counting objects: 593, done.
remote: Compressing objects: 100% (263/263), done.
remote: Total 593 (delta 91), reused 347 (delta 87), pack-reused 240
Receiving objects: 100% (593/593), 8.29 MiB | 11.91 MiB/s, done.
Resolving deltas: 100% (177/177), done.
Checking connectivity... done.
princebirring1992@ubuntu-14:~$ ls
caffe                                cudnn-8.0-linux-x64-v6.0.tgz      demos
cuda                                cudnn-8.0-linux-x64-v6.0.tgz.1    install-14-0
cuda-repo-ubuntu1404-8-0-local-ga2_8.0.61-1_amd64.deb  data                               log
cuda-repo-ubuntu1404-8-0-local-ga2_8.0.61-1_amd64.deb.1  Deep-Learning                     pycharm-comm
princebirring1992@ubuntu-14:~$ cd Deep-Learning/
princebirring1992@ubuntu-14:~/Deep-Learning$ ls
Caffe_  Readme.md  Torch_
princebirring1992@ubuntu-14:~/Deep-Learning$ cd Torch_/
princebirring1992@ubuntu-14:~/Deep-Learning/Torch_$ ls
Demo  GPU_Test  Lua  Mini_Project  Readme.md  Torch  Website_Link
princebirring1992@ubuntu-14:~/Deep-Learning/Torch_$ cd Mini_Project/
princebirring1992@ubuntu-14:~/Deep-Learning/Torch/Mini_Project$ ls
Readme.md  train_mnist.lua  train_mnist_Remove_Bug.lua
princebirring1992@ubuntu-14:~/Deep-Learning/Torch/Mini_Project$
```

- Run the program in ZeroBraneStudio and investigate and Verify its performance.

Step1:

```
princebirring1992@ubuntu-14:~/Deep-Learning/Torch/Mini_Project$ zbstudio
princebirring1992@ubuntu-14:~/Deep-Learning/Torch/Mini_Project$
```

Step2:



Result:

```
2017-11-04 18:14:15 (41.3 MB/s) - 'mnist4.zip' saved [15313661/15313661]
```

```
inflating: ./test.th7
New maxima : 0.106600 @ 1.000000
New maxima : 0.110900 @ 3.000000
Test Accuracy : 0.105200
Program completed in 1330.89 seconds (pid: 2468).
```

Change:

```
local input, target = inputs[idx], targets:narrow(1,idx,1)
```

To:

```
local input, target = inputs[idx], targets[idx]
```

Results:

| Output                                                                                                                                                                                                                                                                                                                                                               | Local console | Markers |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------|---------|
| <pre> New maxima : 0.879700 @ 5.000000 New maxima : 0.890500 @ 7.000000 New maxima : 0.905500 @ 8.000000 New maxima : 0.909300 @ 11.000000 New maxima : 0.911100 @ 12.000000 New maxima : 0.915900 @ 18.000000 New maxima : 0.916300 @ 22.000000 New maxima : 0.920900 @ 23.000000 Test Accuracy : 0.907300 Program completed in 1329.41 seconds (pid: 2977). </pre> |               |         |

#### 4. Modify the program to run on the GPU.

Add the below code in the existing code to run the program on GPU.

```
require 'cunn'

module: cuda()

criterion:cuda()
```

```

require 'dp'
require 'cunn'

-- Load the mnist data set
ds = dp.Mnist()

-- Extract training, validation and test sets
trainInputs = ds:get('train', 'inputs', 'bchw')
trainTargets = ds:get('train', 'targets', 'b')
validInputs = ds:get('valid', 'inputs', 'bchw')
validTargets = ds:get('valid', 'targets', 'b')
testInputs = ds:get('test', 'inputs', 'bchw')
testTargets = ds:get('test', 'targets', 'b')

-- Create a two-layer network
module = nn.Sequential()
module:add(nn.Convert('bchw', 'bf')) -- collapse 3D to 1D
module:add(nn.Linear(1*28*28, 20))
module:add(nn.Tanh())
module:add(nn.Linear(20, 10))
module:add(nn.LogSoftMax())
module:cuda()

-- Use the cross-entropy performance index
criterion = nn.ClassNLLCriterion()
criterion:cuda()

```

Results:

| Output | Local console                                                                                                                                                                                                                                                                                                                                                        | Markers |
|--------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------|
|        | <pre> New maxima : 0.881400 @ 6.000000 New maxima : 0.894200 @ 8.000000 New maxima : 0.894900 @ 9.000000 New maxima : 0.901500 @ 10.000000 New maxima : 0.908900 @ 13.000000 New maxima : 0.911200 @ 16.000000 New maxima : 0.914300 @ 17.000000 New maxima : 0.919200 @ 25.000000 Test Accuracy : 0.921600 Program completed in 1325.70 seconds (pid: 2882). </pre> |         |

- The program is set up to perform stochastic gradient descent. Modify the program to use minibatches. Experiment with different minibatches sizes.

```

1  require 'dpnn'
2  function trainEpoch(module, criterion, inputs, targets, batch_size)
3      for idx=0,(inputs:size(1) - batch_size), batch_size do
4          --local idx = math.random(1,inputs:size(1))
5          local input= inputs[{{idx + 1, idx + batch_size}}]
6
7          local target = targets[{{idx + 1, idx + batch_size}}]
8          --targets:narrow(1,idx,1)
9          -- forward
10         local output = module:forward(input)
11         local loss = criterion:forward(output, target)
12         -- backward
13         local gradOutput = criterion:backward(output, target)
14         module:zeroGradParameters()
15         local gradInput = module:backward(input, gradOutput)
16         -- update
17         module:updateGradParameters(0.9) -- momentum (dpnn)
18         module:updateParameters(0.1) -- W = W - 0.1*dL/dW
19     end
20 end
21
22 bestAccuracy, bestEpoch = 0, 0
23 wait = 0
24 for epoch=1,30 do
25     trainEpoch(module, criterion, trainInputs, trainTargets, 200)
26     local validAccuracy = classEval(module, validInputs, validTargets)
27     if validAccuracy > bestAccuracy then

```

Result:

| Output | Local console                                    | Markers |
|--------|--------------------------------------------------|---------|
|        | New maxima : 0.946700 @ 21.000000                |         |
|        | New maxima : 0.947400 @ 22.000000                |         |
|        | New maxima : 0.947800 @ 23.000000                |         |
|        | New maxima : 0.948100 @ 24.000000                |         |
|        | New maxima : 0.948400 @ 25.000000                |         |
|        | New maxima : 0.948500 @ 26.000000                |         |
|        | New maxima : 0.948600 @ 28.000000                |         |
|        | New maxima : 0.949100 @ 29.000000                |         |
|        | Test Accuracy : 0.951300                         |         |
|        | Program completed in 149.09 seconds (pid: 3038). |         |

| Mini Batches | Timing         |
|--------------|----------------|
| 100          | 103.55 Seconds |
| 200          | 93.43 Seconds  |
| 500          | 95.67 Seconds  |
| 1000         | 96.01 Seconds  |
| 5000         | 93.88 Seconds  |
| 10000        | 98.39 Seconds  |

- Experiment with different numbers of layers and different number of neurons. Maintain the total numbers of weights and biases in the networks, while increasing the number of layers in the network. Describe how the performance changes as the numbers of layers increases – both in terms of training time and performance.

Parameters before adding a layer

| Output                                                                                                                                                                           | Local console | Markers |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------|---------|
| <pre>Program starting as '/opt/zbstudio/bin/linux/x64/lua' -e "io.stc Program 'lua' started in '/opt/zbstudio/myprograms' (pid: 4218). 15910 [torch.LongStorage of size 1]</pre> |               |         |

```
-- Create a two-layer network
module = nn.Sequential()
module:add(nn.Convert('bchw', 'bf')) -- collapse 3D to 1D
module:add(nn.Linear(1*28*28, 19))
module:add(nn.Tanh())
module:add(nn.Linear(19, 33))
module:add(nn.Tanh())
module:add(nn.Linear(33, 10))
module:add(nn.LogSoftMax())
module:cuda()
```

Parameters after adding the layer and different numbers of neurons.

| Output                                                                                                                                                                           | Local console | Markers |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------|---------|
| <pre>Program starting as '/opt/zbstudio/bin/linux/x64/lua' -e "io.stc Program 'lua' started in '/opt/zbstudio/myprograms' (pid: 4228). 15915 [torch.LongStorage of size 1]</pre> |               |         |

- Try one other training function from the list on this page:  
<https://github.com/torch/optim/blob/master/doc/algos.md> . Compare the performance with gradient descent.
- Try different Transfer functions and do comparison study between those. Explain your results.

```

-- Create a two-layer network
module = nn.Sequential()
module:add(nn.Convert('bchw', 'bf')) -- collapse 3D to 1D
module:add(nn.Linear(1*28*28, 19))
module:add(nn.ReLU())
module:add(nn.Linear(19, 33))
module:add(nn.ReLU())
module:add(nn.Linear(33, 10))
module:add(nn.LogSoftMax())
module:cuda()

```

Results:

| Output | Local console                                                                                                                                                                                                                                                                                                                                                                      | Markers |
|--------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------|
|        | New maxima : 0.953600 @ 8.000000<br>New maxima : 0.954200 @ 9.000000<br>New maxima : 0.955500 @ 10.000000<br>New maxima : 0.957400 @ 11.000000<br>New maxima : 0.958900 @ 12.000000<br>New maxima : 0.959400 @ 15.000000<br>New maxima : 0.959700 @ 16.000000<br>New maxima : 0.959800 @ 18.000000<br>Test Accuracy : 0.956600<br>Program completed in 122.21 seconds (pid: 5028). |         |

Execution time increase after changing the transfer function from Tanh() to ReLU().

- Make a chart or table to see the advantages between mini batch, full batch and stochastic gradient. Calculate the timing by using the system clock command.

|                     | Time            | Advantages                                                                                   |
|---------------------|-----------------|----------------------------------------------------------------------------------------------|
| Full Batch = 50,000 | 121.65 Seconds  | Process the entire batch as one as has high learning rate as compare to stochastic gradient. |
| Mini Batch = 10,000 | 108.73 Seconds  | Mini-batch has very high learning rate as compare to Full batch and Stochastic Gradient.     |
| Stochastic Gradient | 1325.70 Seconds |                                                                                              |